

# **Sentiment Analysis Workshop**

**Customer Feedback Classification**

**Using GenAI**

# About Me

- **Current Role:** Learning Experience Manager at [AlignAI](#)



- **Experience:** 15+ years in analytics
- **Passions:** Artificial Intelligence (AI), Governance, and Instructional Design
- **Memberships & Volunteer Work:** WIA and MORPC
- **Education:** Doctoral student at Franklin University



Kristy Wedel

Learning Experience Manager |  
Transforming Data Governance & An...



# Workshop Goals

- Explore AI for sentiment analysis
- Hands-on approach using **Cursor**
- Build, train, and evaluate a sentiment analysis model
- Gain practical skills to adapt and refine models

# Requirements

**Follow Sentiment Analysis Workshop Quick Start.docx Instructions**

## Install Conda

[Miniconda \(lightweight\)](#)

OR

[Anaconda \(comprehensive\)](#)

## Install Cursor

[Download Cursor](#)

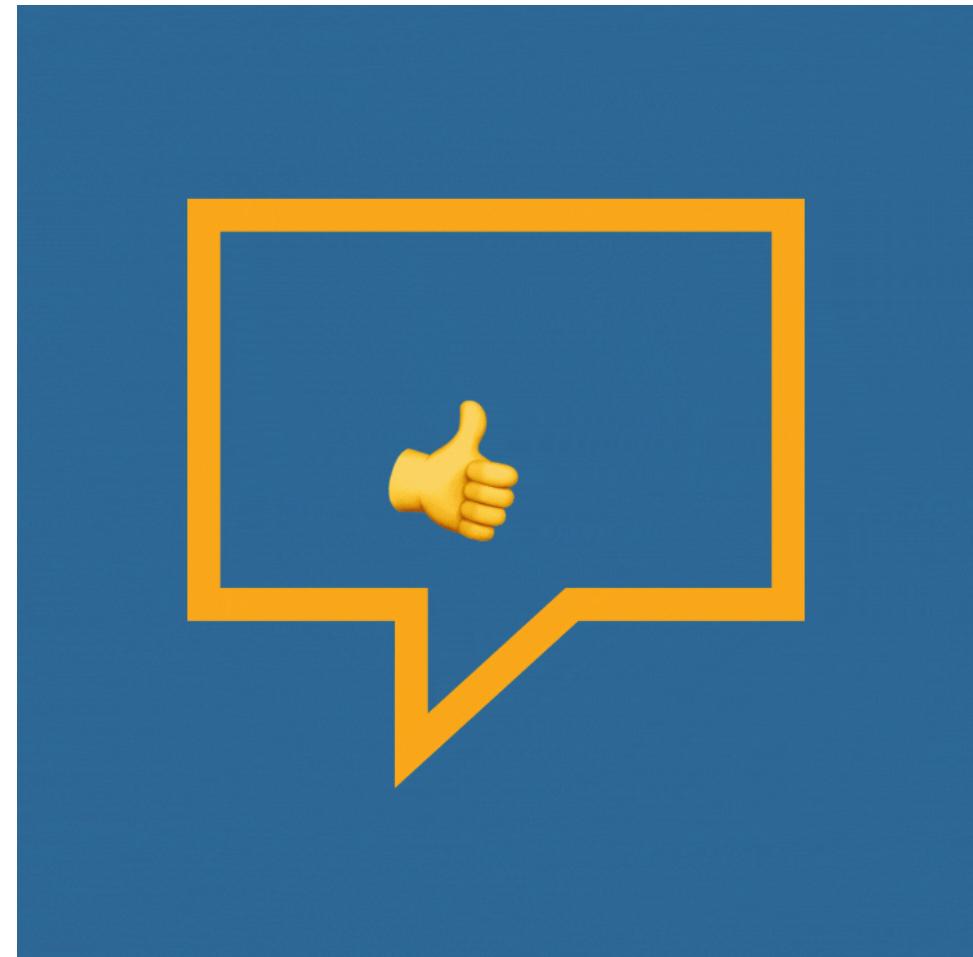
# Files



<https://bit.ly/3PIK4RU>

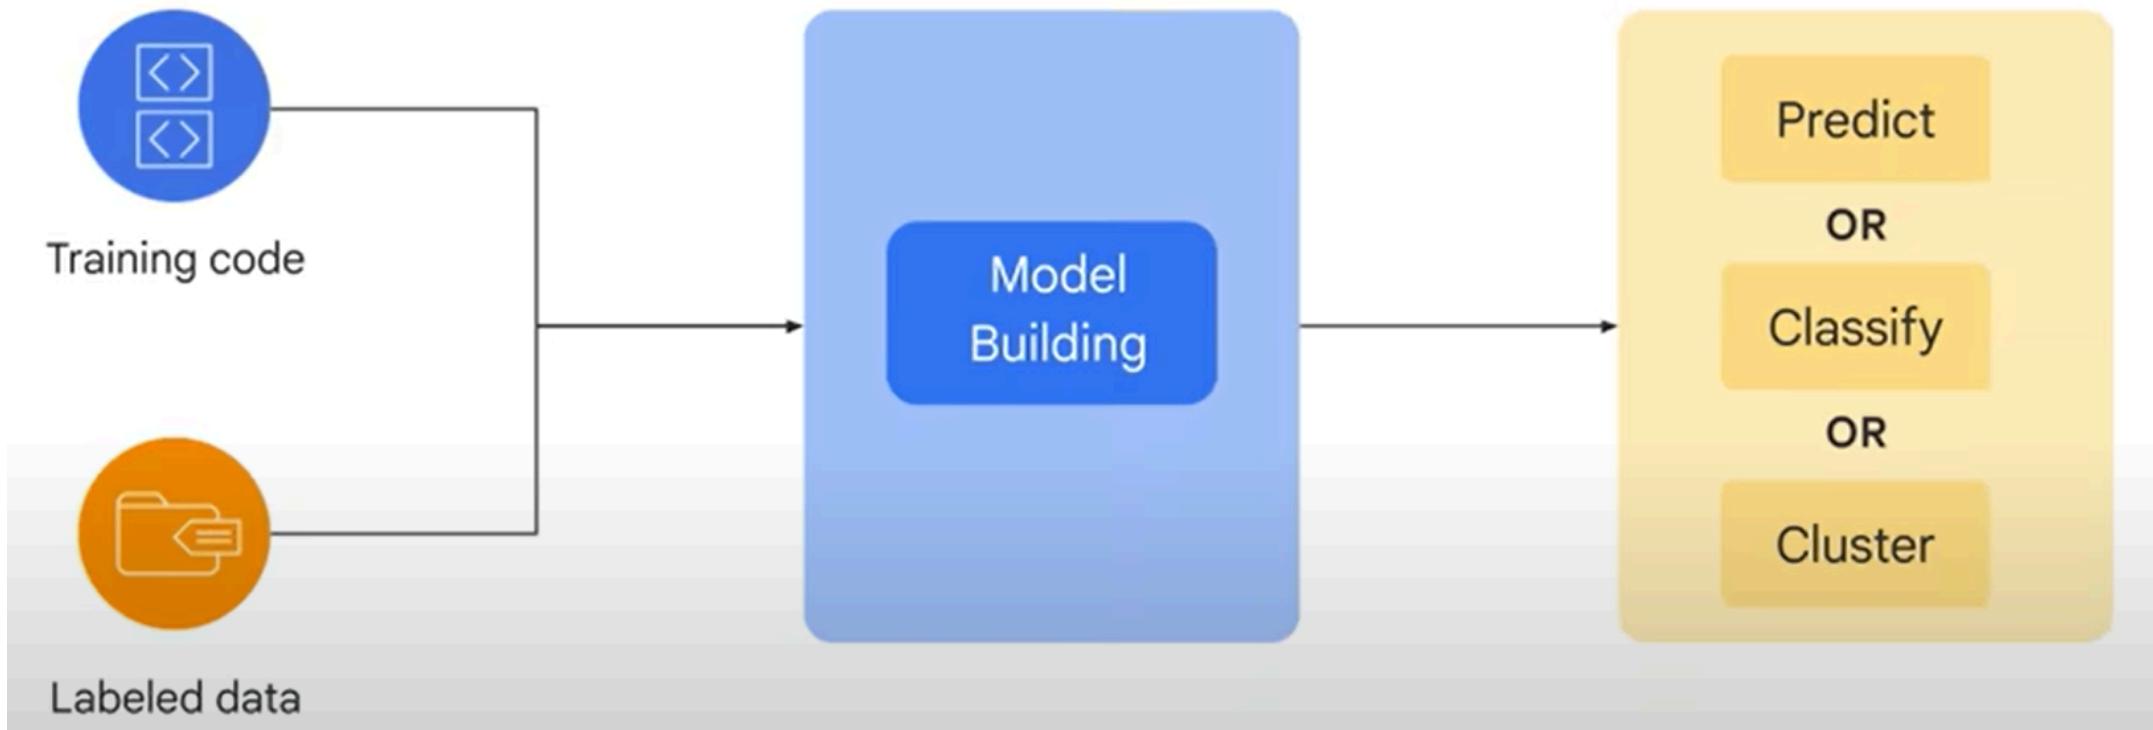
# Introduction to Sentiment Analysis

- NLP technique to classify text into:
  - Positive, Negative, Neutral
- Applications:
  - Customer feedback analysis
  - Social media monitoring
  - Prioritizing support tickets

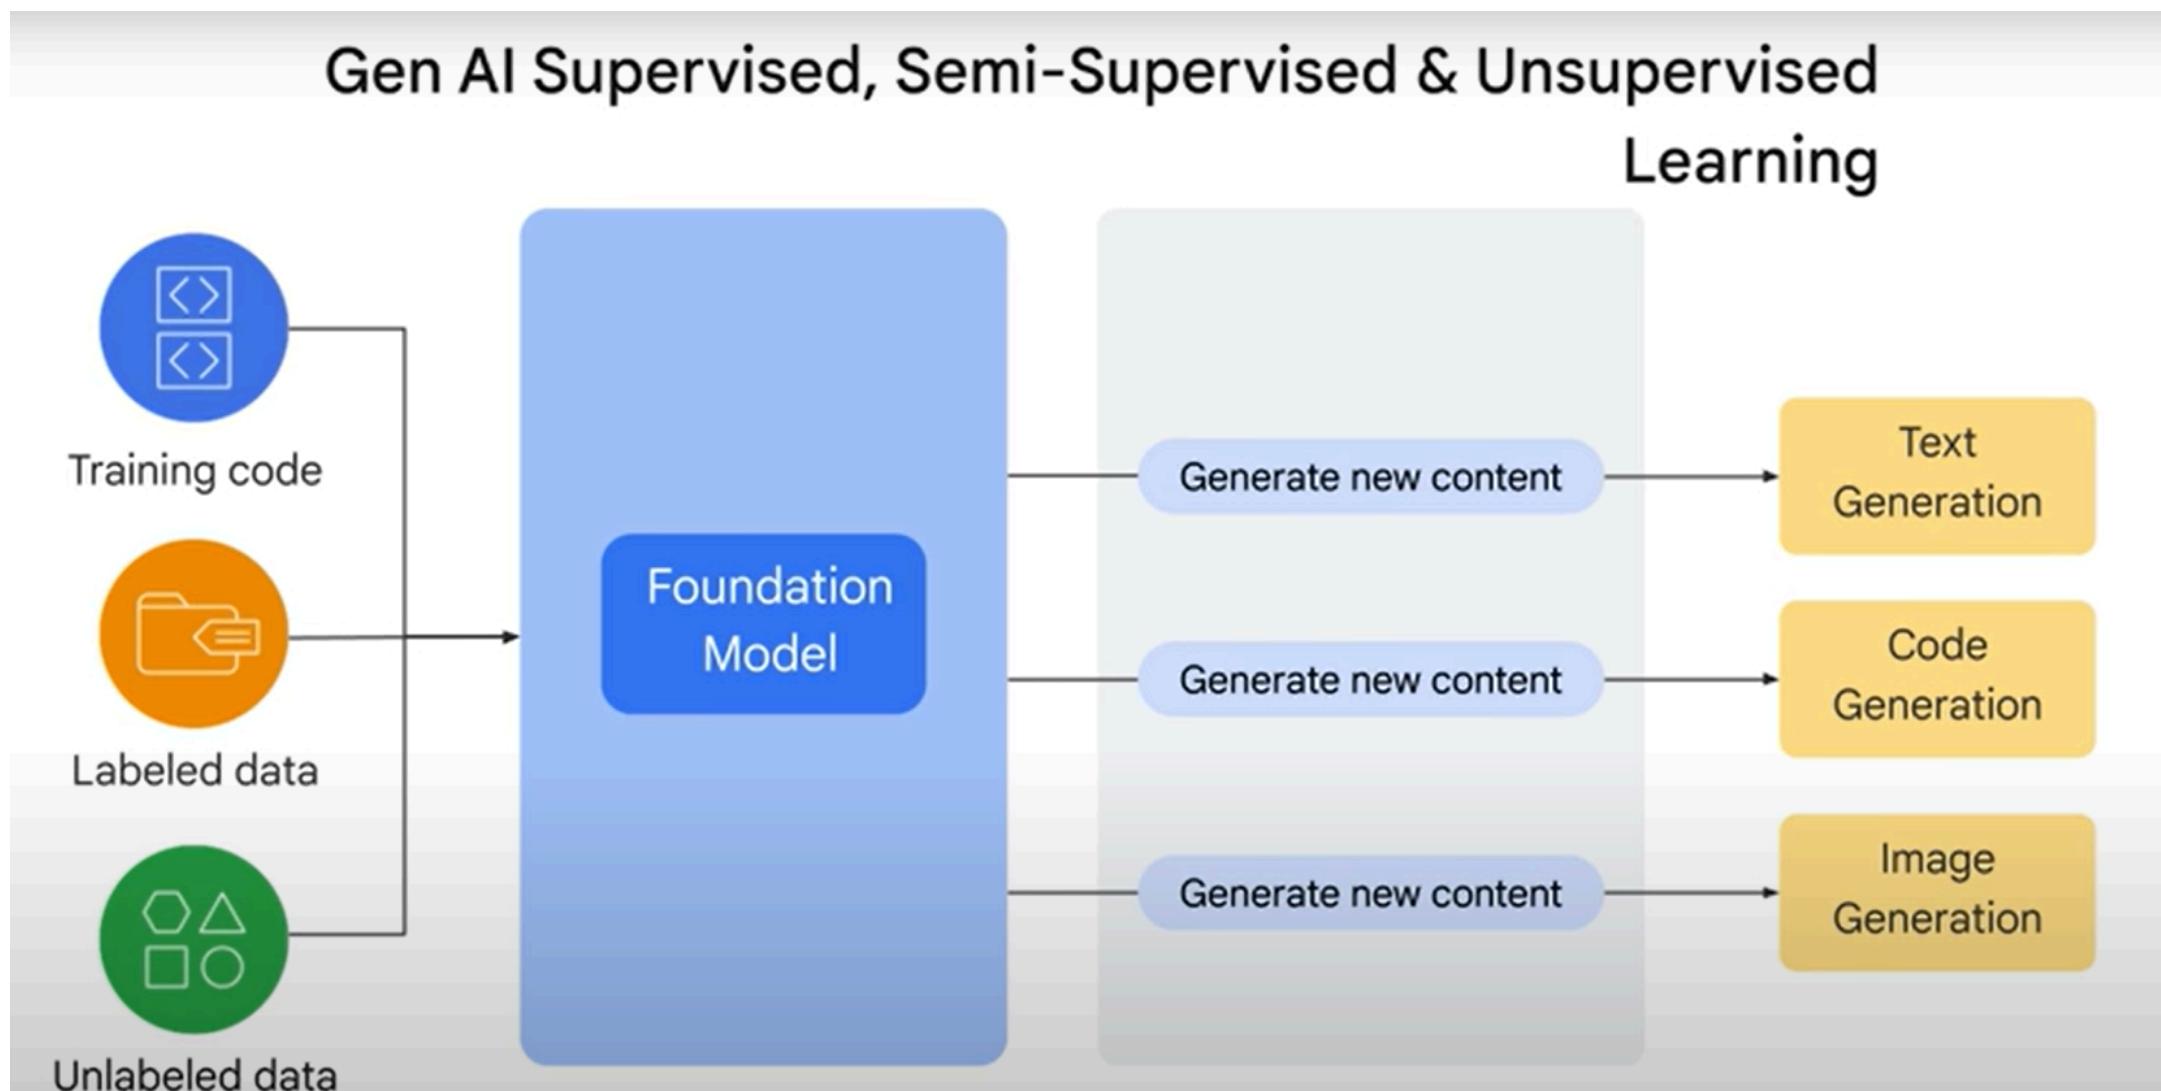


# Traditional ML

## Classical Supervised & Unsupervised Learning



# Gen AI

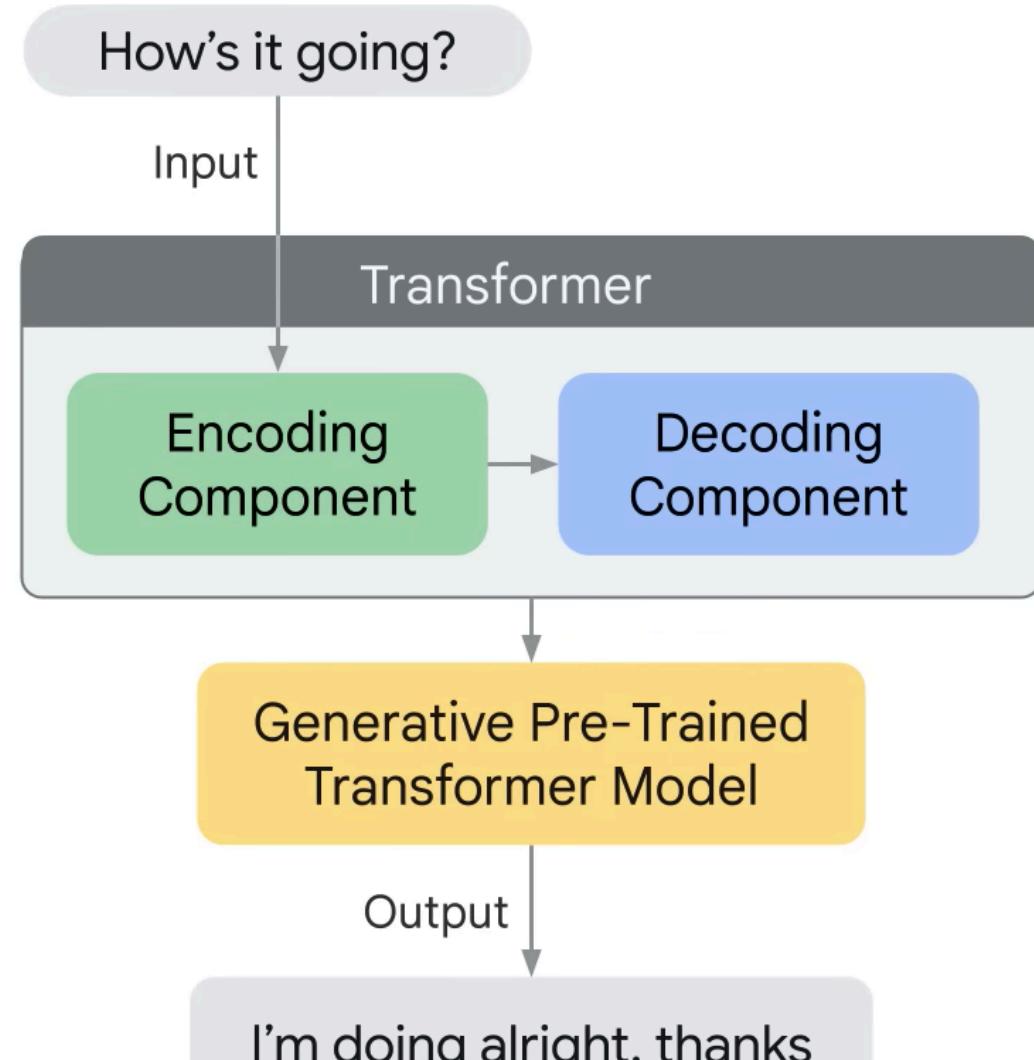


# How Gen AI Works

## How it Works

### Pre-Training:

- Large amount of Data
- Billions of parameters
- Unsupervised learning

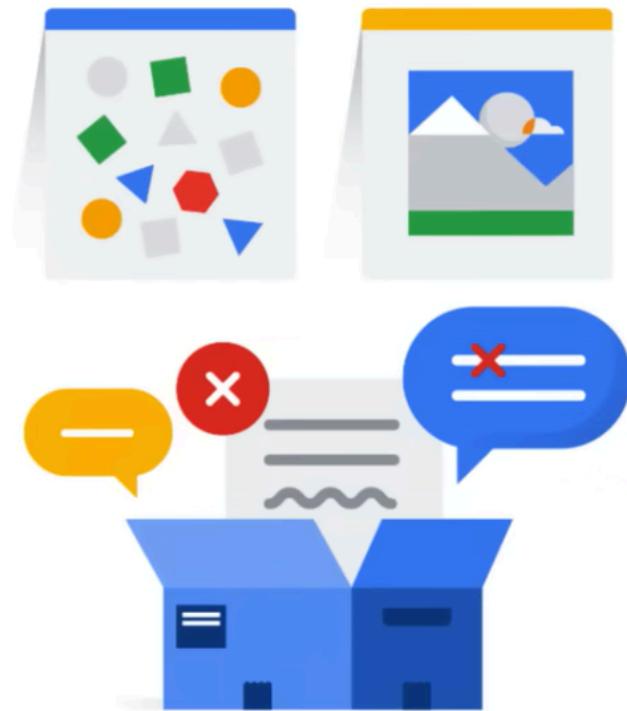


# Hallucinations

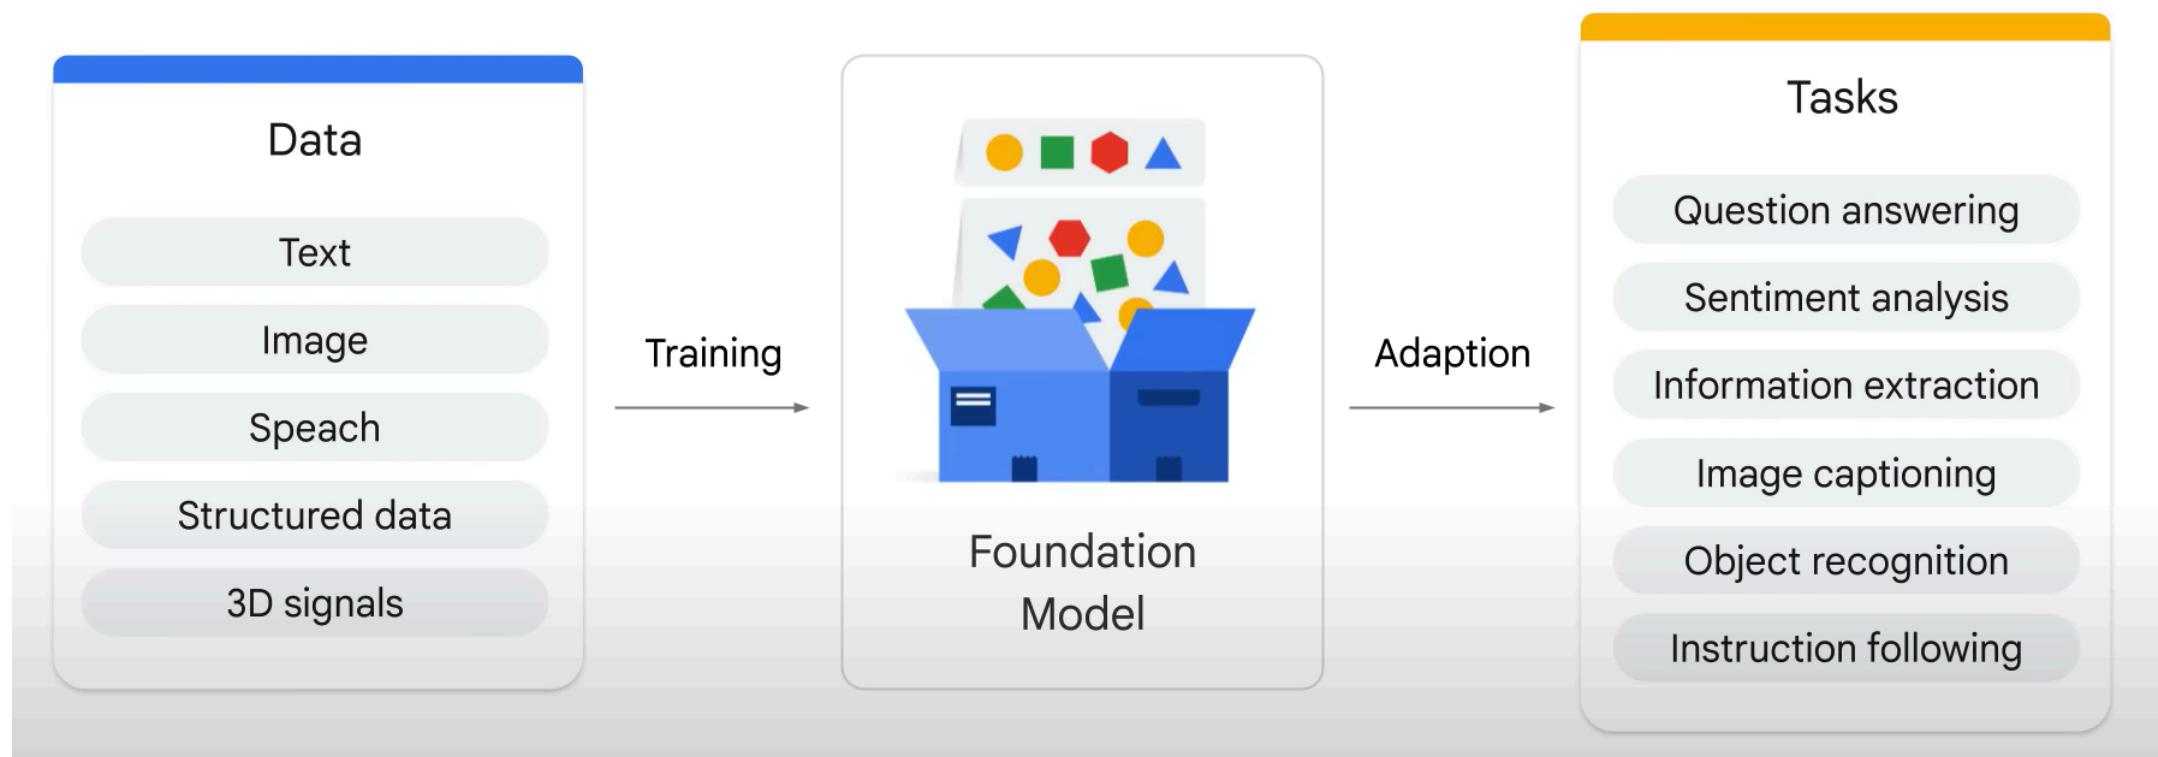
## Hallucinations

### Challenges

- ❗ The model is not trained on enough data
- ❗ The model is trained on noisy or dirty data
- ❗ The model is not given enough context
- ❗ The model is not given enough constraints



# Customization

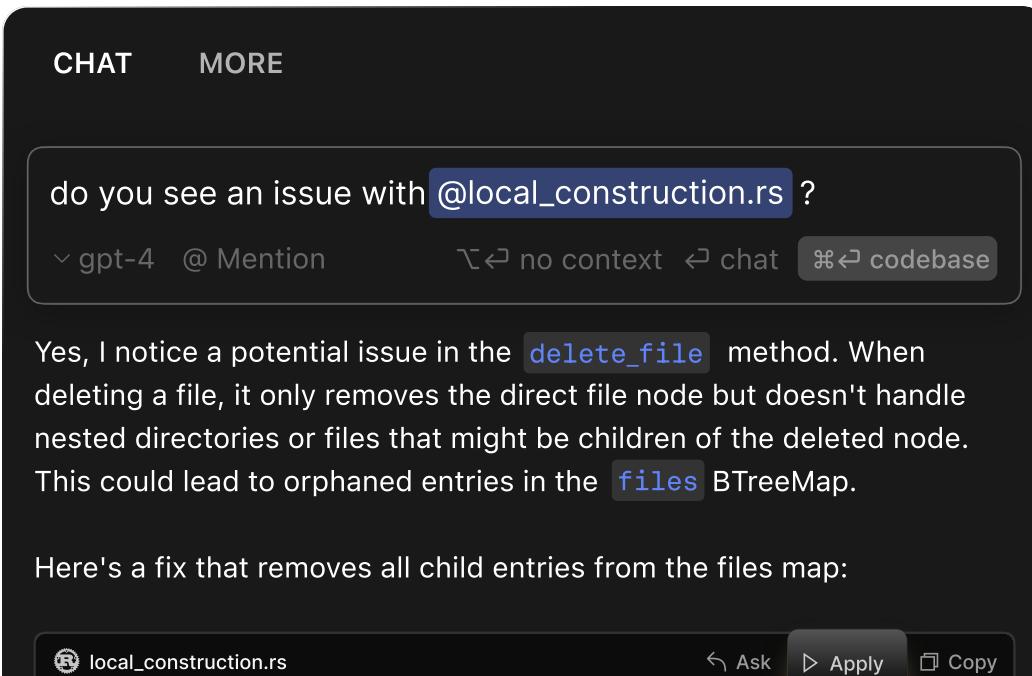


# What is Cursor?

- AI-powered code editor designed to boost developer productivity.
- Combines a clean interface with advanced AI capabilities.

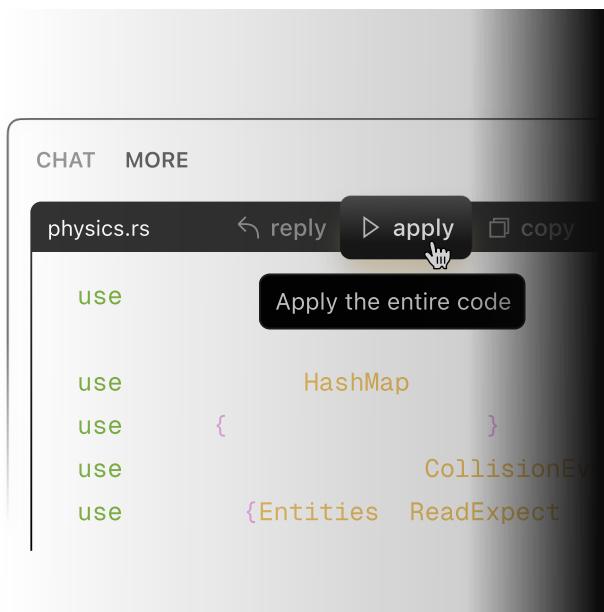
# Cursor: Chat

- Chat - Chat lets you talk with an AI that sees your codebase. The chat can always see your current file and cursor, so you can ask it things like: "Is there a bug here?". You can add particular blocks of code to the context with Ctrl+Shift+L or "@." You can chat with your entire codebase with Ctrl+Enter.



# Cursor: Instant Apply

- Instant Apply - Apply the code suggestions from chat back into your codebase by clicking the play button on top of any chat codeblock.



# Cursor Benefits

- Saves time by reducing repetitive tasks.
- Enhances code quality with AI recommendations.
- Simplifies onboarding for new developers.
- Improves collaboration with integrated tools.

# Workflow Steps

**Step 1:** Data Cleaning and Preprocessing

**Step 2:** Model Training and Evaluation

**Step 3:** Deployment and Prediction

**Step 4:** Real-Time Sentiment Predictions

# Data Overview

- **Dataset Details:**
  - 1,000 customer feedback entries.
  - Columns: Review\_ID , Review\_Text , Sentiment .
  - Labels: Positive, Neutral, Negative.

# Google Colab

<https://tinyurl.com/3jccy5yc>

# Workflow Steps

**Step 1:** Data Cleaning and Preprocessing

**Step 2:** Model Training and Evaluation

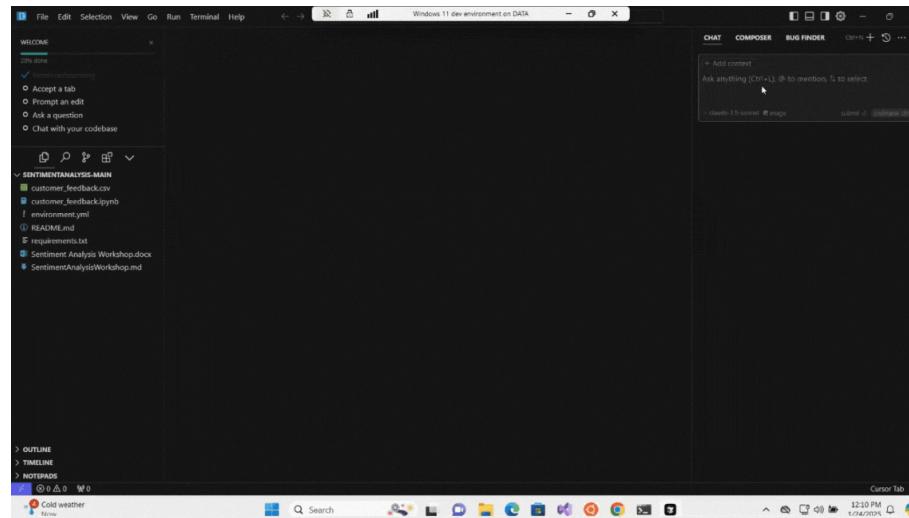
**Step 3:** Deployment and Prediction

# Load Data

## Prompt Example:

Write Python code to:

- Create a python jupyter notebook to load customer\_feedback.csv



# Data Preprocessing

- Data preprocessing is a critical step in data analysis and machine learning.
- It ensures the raw data is clean, structured, and usable for analysis.

## Key Steps

- **Data Cleaning:** Handle missing values, remove duplicates, and correct inconsistencies.
- **Data Transformation:** Normalize, encode variables, and create new features.
- **Data Integration:** Merge data from multiple sources, resolve schema mismatches.
- **Data Reduction:** Reduce dataset size while preserving important information.

# Cleaning for Sentiment Analysis

- **Text Cleaning:**
  - Remove special characters and stop words and convert to lowercase.
- **Tokenization:**
  - Split `Review_Text` into words for analysis.
- **Lemmatization:**
  - Reduce words to their base forms (e.g., "running" → "run").
- **Handling Missing Data:**
  - Drop or impute missing entries.
- **Deduplication:**
  - Remove duplicate reviews to avoid bias.

# Data Cleaning

## Prompt Example:

Write Python code to replace the existing data:

- Remove duplicate rows.
- Drop rows where "Review\_Text" or "Sentiment" columns have missing values.

Then, save the dataset as a new csv and visualize the distribution by sentiment with a simple chart.

The screenshot shows a Jupyter Notebook environment on a Windows desktop. The notebook has a single cell containing Python code. The code imports pandas, numpy, and matplotlib, reads a CSV file, displays its first 5 rows, and then performs several operations: it checks for missing values in the dataset, prints the count of missing values, and then drops rows with missing values from the DataFrame. A tooltip on the right provides instructions for using the notebook, such as loading the CSV file and running each cell sequentially.

```
# Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Read the CSV file
df = pd.read_csv('customer_feedback.csv')

# Display the first 5 rows of the dataset
print("First 5 rows of the dataset:")
display(df.head())

# Get basic information about the dataset
print("Basic information about the dataset")
print(df.info())

# Display basic statistics of numerical columns
print("Basic statistics of numerical columns")
display(df.describe())

# Check for missing values
print("Missing values:")
print(df.isnull().sum())

... Ahs ...
```

Review_ID	Review_Text	Sentiment	
0	522	Excellent service and fast delivery. Highly recomended.	Positive
1	728	Terrible experience. The product broke after a few days.	Negative
2	741	Neutral feelings about this product. It's functional.	Neutral
3	661	This is the best purchase I've made this year!	Positive
4	412	The product is okay, but nothing special.	Neutral

# Workflow Steps

**Step 1:** Data Cleaning and Preprocessing

**Step 2:** Model Training and Evaluation

**Step 3:** Deployment and Prediction

# Generative AI Key Terms

- **Tokens:** The smallest units of text, such as words or subwords, used in AI processing.
- **Encoding:** The process of converting text into numerical representations for machine learning models.
- **Transformer:** A neural network architecture that uses self-attention to process sequences of data.
- **Loss:** A metric that measures the difference between predicted and actual outputs, guiding model training.
- **Epochs:** Complete passes through the entire training dataset during model training.

# Overview of Pretrained Models

## **distilbert-base-uncased-finetuned-sst-2-english**

- **Base Model:** DistilBERT, a lighter, faster version of BERT.
- **Fine-Tuned Task:** Sentiment analysis using the SST-2 dataset.
- **Performance:** Delivers accurate binary sentiment predictions (positive/negative).
- **Use Cases:** Ideal for analyzing customer feedback, social media, and reviews.

# Why Use Pretrained Models?

- Minimize training time and computational resources.
- Achieve state-of-the-art results with minimal fine-tuning.
- Versatile across various NLP tasks.

# Key Preprocessing Steps for DistilBERT

Step	Example	Explanation
Convert to Lowercase	"Great Day" → "great day"	Ensures uniformity by treating uppercase and lowercase words as the same.
Remove Punctuation	"Hello, world!" → "Hello world"	Eliminates unnecessary symbols to focus on meaningful content.
Remove Stop Words	"This is an example" → "example"	Removes common words that do not contribute to understanding the meaning.
Apply Lemmatization	"running", "runs" → "run"	Reduces words to their base forms, grouping similar words together.
Tokenization	"I love AI" → ["I", "love", "AI"]	Splits text into smaller units (tokens) for easier processing.
Padding/Truncation	["I", "love"] → ["I", "love", "[PAD]"]	Ensures all inputs are of uniform length for batch processing.
Add Special Tokens	["I", "love"] → "[CLS]", "I", "love", "[SEP]"	Special tokens [CLS] and [SEP] mark the start and end of the input.
Create Attention Masks	["I", "love", "[PAD]"] → [1, 1, 0]	Binary masks differentiate real input tokens from padding.
Numerical Encoding	["I", "love"] → [101, 2027]	Converts tokens into numerical IDs for processing by the model.

# Other Pretrained Models

## 1. BERT:

- Bidirectional context understanding for tasks like Q&A and classification.
- Popular variants: BERT-Base, BERT-Large.

## 2. GPT:

- Autoregressive model for text generation and conversational AI.

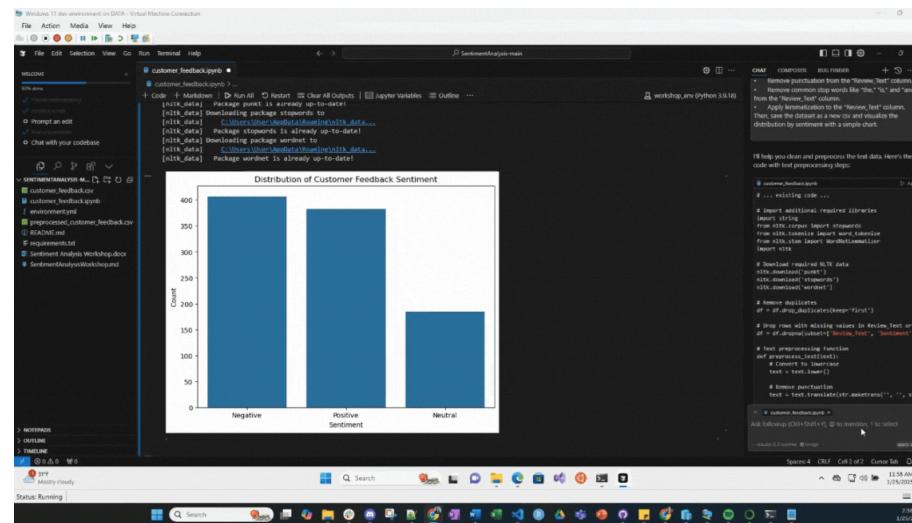
## 3. DistilBERT:

- Lightweight, fast, and efficient for edge devices.

## Prompt Example:

Use Hugging Face Transformers and tensorflow to train a sentiment classification model:

- Model: distilbert-base-uncased-finetuned-sst-2-english.
- Group Positive and Neutral sentiments together.
- Provide metrics to evaluate: Accuracy, Precision, Recall, F1-score.



# Overfitting vs. Underfitting in Machine Learning Models

# What is Overfitting?

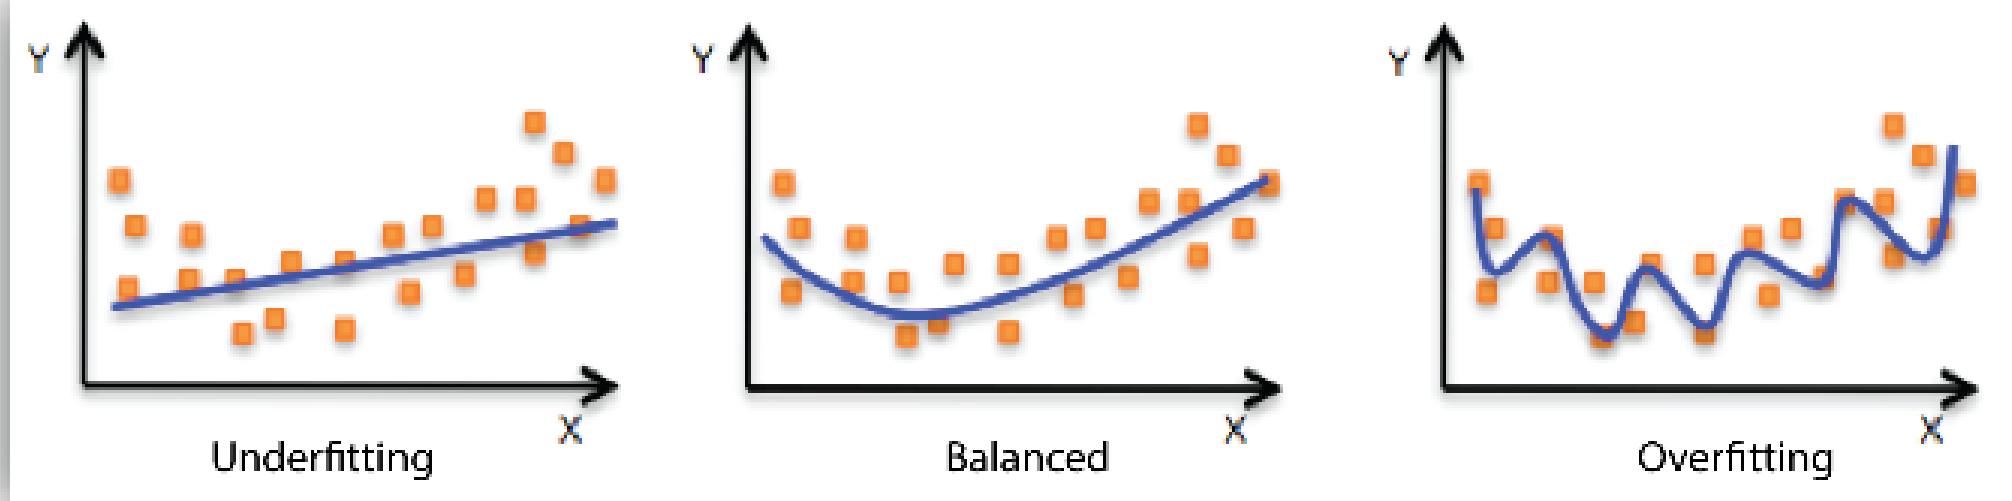
- **Definition:** When a model performs well on training data but poorly on unseen data.
- **Indicators:**
  - High training accuracy, low validation/test accuracy.
  - Large gap between training and validation loss.
- **Causes:**
  - Model memorizes training data instead of generalizing patterns.
  - Excessive model complexity (e.g., too many layers or parameters).

# What is Underfitting?

- **Definition:** When a model performs poorly on both training and unseen data.
- **Indicators:**
  - Low accuracy on training data.
  - Training and validation loss remain high.
- **Causes:**
  - Model fails to capture underlying patterns in the data.
  - Model is too simple (e.g., insufficient layers or parameters).

# Key Comparison

Overfitting	Underfitting
Memorizes training data	Fails to learn patterns
High training accuracy, low test accuracy	Low accuracy overall
Caused by excessive complexity	Caused by insufficient complexity



# **Strategies to Address**

## **1. For Overfitting:**

- Data augmentation, dropout, reduce model complexity, K-fold cross-validation.

## **2. For Underfitting:**

- Use a more complex model, increase training time, improve feature engineering, or collect more data.

# Workflow Steps

**Step 1:** Data Cleaning and Preprocessing

**Step 2:** Model Training and Evaluation

**Step 3:** Deployment and Prediction

# Prompt Example:

Deploy a simple FastAPI application for sentiment analysis with:

1. TensorFlow integration.
2. Environment variables to suppress TensorFlow warnings.
3. Endpoint: "/predict" to classify sentiment.
4. An application with an app.py and app\_test.py. The app\_test.py file should have a visual output.
5. There should also be an interactive front end.

The screenshot shows a Jupyter Notebook interface running on a Windows 11 desktop. The notebook title is "SentimentAnalysis-main". The code cell contains Python code for initializing a TensorFlow model and performing sequence classification. The output cell shows training logs, a classification report, and a line chart titled "Model Accuracy over Epochs".

**Classification Report:**

	precision	recall	f1-score	support
Negative	1.00	1.00	1.00	81
Positive/Neutral	1.00	1.00	1.00	114
accuracy	1.00	1.00	1.00	195
macro avg	1.00	1.00	1.00	195
weighted avg	1.00	1.00	1.00	195

**Model Accuracy over Epochs:**

The chart shows a single blue line representing the model's accuracy over 13 epochs. The accuracy starts at approximately 0.99 and remains stable throughout the training process.

# Conclusion

- **Key Steps:**

- Data preparation.
- Model training and evaluation.
- FastAPI deployment.

- **Takeaways:**

- High-quality data is critical.
- Deployment enables real-time applications.
- Simplifying sentiment categories impacts granularity.

# Thank You!

- Slides are located in **SentimentAnalysisWorkshop.pdf**
- A completed project is located in **CompletedProject.zip**



Kristy Wedel  
Learning Experience Manager |  
Transforming Data Governance & An...

